

Network Forensics and Analysis Poster

Continuous Incident Response and Threat Hunting: Proactive Threat Identification

CORE CONCEPT:

Apply new intelligence to existing data to discover unknown incidents

NETWORK FORENSICS USE CASE:

Threat intelligence often contains network-based indicators such as IP addresses, domain names, signatures, URLs, and more. When these are known, existing data stores can be reviewed to determine if there were indications of the intel-informed activity that warrant further investigation.

Post-Incident Forensic Analysis: Reactive Detection and Response

CORE CONCEPT:

Examine existing data to more fully understand a known incident

NETWORK FORENSICS USE CASE:

Nearly every phase of an attack can include network activity. Understanding an attacker's actions during Reconnaissance, Delivery, Exploitation, Installation, Command and Control, and Post-Exploitation phases can provide deep and valuable insight into their actions, intent, and capability.

Network Forensics is a critical component for most modern digital forensic, incident response, and threat hunting work. Whether pursued alone or as a supplement or driver to traditional endpoint investigations, network data can provide decisive insight into the human or automated communications within a compromised environment.

Network Forensic Analysis techniques can be used in a traditional forensic capacity as well as for continuous incident response/threat hunting operations.

Additional Resources

SANS FOR572: Advanced Network Forensics and Analysis: for572.com/course

FOR572 Course Notebook: for572.com/notebook

Network Forensics and Analysis Poster: for572.com/poster

GIAC Certified Network Forensic Analyst certification available: gnfa.org

SOF-ELK®



What is "ELK" and the "Elastic Stack"?

The Elastic Stack consists of the Elasticsearch search and analytics engine, the Logstash data collection and enrichment platform, and the Kibana visualization layer. It is commonly known as "ELK", named for these three components.

The broader Elastic Stack includes other components such as the Elastic Beats family of log shippers, and various security and performance monitoring components.

All of the ELK components and the Beats log shippers are free and open-source software. Some other components of the Elastic Stack are commercially-licensed.

Booting and Logging into SOF-ELK

The SOF-ELK VM is distributed in ready-to-boot mode. You may want to add additional CPU cores and RAM if available. Do not decrease the CPU or RAM. After the VM boots, its IP address is displayed on the pre-authentication screen. This IP address is needed for both remote shell access (SSH) and web access to the Kibana interface. Log in with the "elk_user" account and password "forensics". The elk user has administrative access using the sudo utility. The password should be changed after first login using local preferences or policies. The SSH server is running on the default port, 22. Access this with your preferred SSH/SCP/SCP/FTP client software. The Kibana interface is running on port 5601. Access this with your preferred web browser. (Note that Microsoft Edge is known to be problematic.)

Kibana Query Language Syntax

The Kibana user interface uses the Kibana Query Language (KQL) syntax for searching the data contained in Elasticsearch. Below are some of the basic syntaxes that will help you search data that has been loaded to SOF-ELK.

Basic Searching

The most basic search syntax is "fieldname:value", which will match all documents with a "fieldname" field set to a value of "value". Searches can be negated by prefixing them with "not". Some examples:

- hostname:webserver
- not querystype:AAAA

Logical Construction

Multiple searches can be combined using "and" and "or".

- destination_geo.asn:Amazon.com and in_bytes > 1000000

Numerical Ranges

Fields containing numerical values can be searched with standard range operators.

- total_bytes > 1000000
- return_code > 200 and return_code < 300

Partial String Searches

The "*" is used as a wildcard character.

- username:admin*
- query:*.cz.cc

IP Addresses and CIDR Blocks

IP address fields can be searched for specific values or use CIDR notation for netblocks.

- source_ip:172.16.7.11
- destination_ip:172.16.6.0/24

SOF-ELK is a VM appliance with a preconfigured, customized installation of the Elastic Stack. It was designed specifically to address the ever-growing volume of data involved in a typical investigation, as well as to support both threat hunting and security operations components of information security programs. The SOF-ELK customizations include numerous log parsers, enrichments, and related configurations that aim to make the platform a ready-to-use analysis appliance. The SOF-ELK platform is a free and open-source appliance, available for anyone to download. The configuration files are publicly available in a GitHub repository and the appliance is designed for upgrades in the field. The latest downloadable appliance details are at for572.com/sof-elk-readme.

Loading Data to SOF-ELK

SOF-ELK can ingest several data formats, including:

- Syslog (many different log types supported)
- NetFlow
- Selected Zeek logs
- HTTP server access logs
- EZ Tools JSON files

More sources are being tested and added to the platform and can be activated through the GitHub repository. See the "Updating With Git" section for more details on how to do this.

All source data can be loaded from existing files (DFIR Model) as well as from live sources (Security Operations Model).

DFIR Model

Place source data onto the SOF-ELK VM under the /logstash/ directory tree.

Syslog data: /logstash/syslog/

Since syslog entries often do not include the year, subdirectories for each year can be created in this location – for example,

/logstash/syslog/2018/

HTTP Server logs: /logstash/httpd/

Supports common, combined, and related formats

PassiveDNS logs: /logstash/passivedns/

Raw logs from the passivedns utility

NetFlow from nfcapd_collected data stores:

/logstash/nfarch/

Use the included nfcdump2sof-elk.sh or vpcflow2sof-elk.sh scripts to create SOF-ELK-compatible NetFlow ASCII files.

Zeek NSM logs: /logstash/zeek/

Supports multiple different log types, based on default Zeek NSM filenames

EZ Tools JSON Files: /logstash/kafe/

Supports multiple files from the KAFE family of Eric Zimmerman's tools in JSON format. Open the necessary firewall port(s) to allow your preferred network-based ingest to occur.

Querying Available Data

The top of each dashboard allows the user to input KQL queries, detailed in the "Kibana Query" section. Elasticsearch determines how well its documents match, including a ".score" field that indicates how well each document matches the query.

Filtering

Filters can also be applied in the Kibana interface. These are similar to queries, but are a binary match/no-match search without a ".score" field. Elasticsearch caches frequently-used filters to optimize their performance.

Kibana shows filters as boxes below the query field. "Must have" and "must not have" are differentiated with the red "NOT" text.

Filters can be modified with the drop-down menu displayed after clicking on a filter.

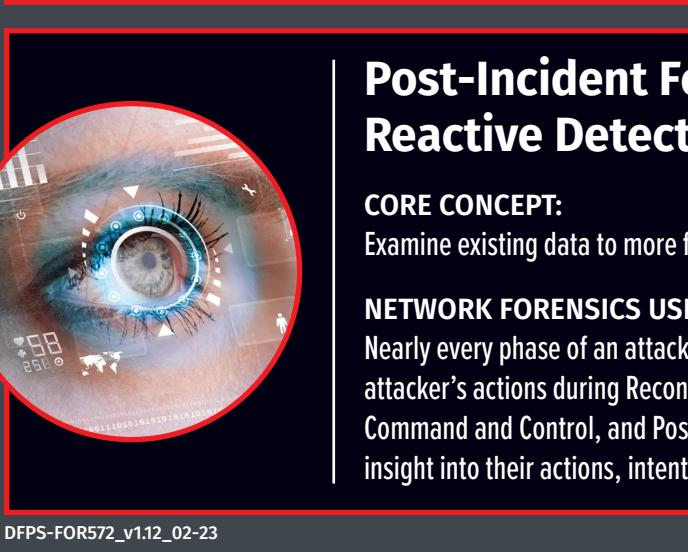
Document Expansion

When a dashboard includes a document listing panel, each document can be expanded by clicking the triangle icon on the left.

This will show all fields for the document.

Interactive Filter Generation

Each field displayed in the record details can be interactively built into a filter with the magnifying glass icons displayed when hovering over the field. The plus sign magnifying glass creates a "must have" filter, the minus sign magnifying glass creates a "must not have" filter. The table icon adds the field to the document listing panel and the final icon creates a "field must be present" filter.



Network Source Data Types



Full-Packet Capture (pcap)

pcap files contain original packet data as seen at the collection point. They can contain partial or complete packet data.

Benefits

- Often considered the "holy grail" of network data collection, this data source facilitates deep analysis long after the communication has ended.
- Countless tools can read from and write to pcap files, giving the analyst many approaches to examine them and extract relevant information from them.

Drawbacks

- These files can grow extremely large – tens of terabytes of pcap data can be collected each day from a 1Gbps link. This scale often makes analysis challenging.
- Legal constraints often limit availability of this source data. Such constraints are also complicated when an organization crosses legal jurisdictions.
- Encrypted communications are increasingly used, rendering full-packet capture less useful for low-level analysis.



NetFlow and Related Flow-Based Collections

Flow records contain a summarization of network communications seen at the collection point. NetFlow contains no content – just a summary record including metadata about each network connection. Whether used alone to determine if communications occurred or in conjunction with other data sources, NetFlow can be extremely helpful for timely analysis.

Benefits

- Since they are collected and retained for business operations purposes, logs are widely available and processes often in place to analyze them.
- Raw log data can be aggregated for centralized analysis. Many organizations have this capability in some form of SIEM or related platform.

Drawbacks

- Log data contains varying levels of detail in numerous formats, often requiring parsing and enrichment to add context or additional data to corroborate findings.
- If log data is not already aggregated, finding it can involve significant time and effort before analysis can begin.



Log Files

Log files are perhaps the most widely-used source data for network and endpoint investigations. They contain application or platform-centric items of use to characterize activities handled or observed by the log creator.

Benefits

- Since they are collected and retained for business operations purposes, logs are widely available and processes often in place to analyze them.
- Analysis processes are much faster with NetFlow than full-packet capture. It can be 100-1000x faster to run a query against NetFlow than the corresponding pcap file.
- There are generally fewer privacy concerns with collecting and storing NetFlow. Local legal authority should be consulted prior to use.
- Analysis processes apply equally to all protocols – encrypted or plaintext, custom or standards-based.

Drawbacks

- Log data is often aggregated, finding it can involve significant time and effort before analysis can begin.
- If log data is not already aggregated, finding it can involve significant time and effort before analysis can begin.

Network-Based Processing Workflows

Although there is no single workflow to exhaustively perform network forensic analysis, the most common and beneficial tasks can generally be placed into the categories below. Note that these categories are not generally iterative. They are components of a dynamic process that can adapt to adversaries' actions.

Ingest and Distill

GOAL: Prepare for analysis and derive data that will more easily facilitate the rest of the analytic workflow

- Log source data according to local procedure
- If pcap files are available, distill to other data source types (NetFlow, Zeek logs, Passive DNS logs, etc.)
- Consider splitting source data into time-based chunks if the original source covers an extended period of time
- Load source data to large-scale analytic platforms such as SOF-ELK, Arkime, etc.

Reduce and Filter

GOAL: Reduce large input data volume to a smaller volume, allowing analysis with a wider range of tools

- Reduce source data to a more manageable volume using known indicators and data points
- Initial indicators and data points may include IP addresses, ports/protocols, time frames, volume calculations, domain names and hostnames, etc.
- For large-scale analytic platforms, build filters to reduce visible data to traffic involving known indicators

Extract Indicators and Objects

GOAL: Find artifacts that help identify malicious activity, including field values, byte sequences, files, or other objects

- As additional artifacts are identified, maintain an ongoing collection of these data points for further use during and after the investigation
- These may include direct observations from within the network traffic or ancillary observations about the nature of the communications – related DNS activity, before/after events, etc.
- Extracting files and other objects such as certificates or payloads can help feed other parts of the IR process such as malware reverse engineering and host-based activity searches
- Protect this data according to local policies and share in accordance with appropriate operational security constraints

Scope and Scale

GOAL: Search more broadly within source data for behavior that matches known indicators

- After identifying useful artifacts that define activity of interest, scale up the search using large-scale analytic platforms and tools
- Identify additional endpoints that exhibit the suspicious behavior, aiming to fully scope the incident within the environment
- Pass appropriate indicators to security operations for live identification of suspicious activity

Analyze and Explore

GOAL: Identify traffic and artifacts that support investigative goals and hypotheses

- Within the reduced data set, seek knowledge about the suspicious traffic
- This may include evaluating traffic contents, context, anomalies, consistencies – anything that helps to clarify its relevance to the investigation

• Seek any protocol anomalies that could indicate traffic being misused for suspicious purposes

• Use any available environmental baselines to identify deviations from normal traffic behaviors

Establish Baselines

GOAL: Identify parameters for "normal" patterns of behavior to help find anomalies that need to be investigated

- Determine typical cycles of traffic, top-talking hosts, ports/protocols, GET/POST ratio for HTTP activity, etc.
- Build all baselines for multiple periods – metrics must have different cycles for daily, weekly, monthly, and annual time frames
- Consider the levels within the organization at which the baselines should be built – enterprise-level rollups will generally differ from those at lower levels

Distilling Full-Packet Capture Source Data



While full-packet capture is often collected strategically as a component of a continuous monitoring program or tactically during incident response actions, it is often too large to process n

Network Forensic Toolbox

Tools are a critical part of any forensic process, but they alone cannot solve problems or generate findings. The analyst must understand the available tools and their strengths and weaknesses, then assess the best approach between raw source data and the investigative goals at hand. The tools detailed here are far from a comprehensive list, but represent a core set of utilities often used in network forensic analysis. More extensive documentation is available in the tools' man pages and online documentation.



tcpdump: Log or parse network traffic
Classically used to dump live network traffic to pcap files, **tcpdump** is more commonly used in network forensics to perform data reduction by reading from an existing pcap file, applying a filter, then writing the reduced data to a new pcap file. **tcpdump** uses the BPF (Berkeley Packet Filter) language for packet selection.

Usage:
\$ tcpdump <options> <bpf filter>

Common command-line parameters:

- n Prevent DNS lookups on IP addresses. Use twice to also prevent port-to-service lookups
- r Read from specified pcap file instead of the network
- w Write packet data to a file
- i Specify the network interface on which to capture
- s Number of bytes per packet to capture
- C Maximum of megabytes to save in a capture file before starting a new file
- G Number of seconds to save in each capture file (requires time format in output filename)
- W Used with the -C option, limit the number of rotated files

Note: The BPF filter is an optional parameter

Common BPF primitives:

```
host IP address or FQDN      top Layer 4 protocol is TCP
net Netblock in CIDR notation  udp Layer 4 protocol is UDP
port TCP or UDP port number   icmp Layer 4 protocol is ICMP
ip Layer 3 protocol is IP
```

Parameters such as host, net, and port can be applied in just one direction with the **src** or **dst** modifiers. Primitives can be combined with **and**, **or**, or **not**, and/or can be enforced with parentheses.

BPF Examples:

- * proto tcp and port 80
- * udp and dst host 8.8.8.8
- * src host 1.2.3.4 and (dst net 10.0.0.0/8 or dst net 172.16.0.0/12)

Capturing live traffic generally requires elevated operating system permissions (e.g. **sudo**), but reading from existing pcap files only requires filesystem-level read permissions to the source file itself.

Examples:

```
$ tcpdump -n -r infile.pcap -w -t80.pcap
$ sudo tcpdump -n -i emps3 -w outfile.pcap
$ sudo tcpdump -n -i emps3 -C 1024 -G 100 -w 10GB_rolling_buffer.pcap
$ sudo tcpdump -n -i emps3 -G 86400 -w dns-%T.pcap
```

editcap: Modify contents of a capture file
Since the BPF is limited to evaluating packet content data, a different utility is required to filter on pcap metadata. This command will read capture files, limit the time frame, file size, and other parameters, then write the resulting data to a new capture file, optionally de-duplicating packet data.

Usage:
\$ editcap <options> <input file> <output file>

Common command-line parameters:

- A Select packets at or after the specified time (Use format: YYYY-MM-DD HH:MM:SS)
- B Select packets before the specified time
- d De-duplicate packets (Can also use -D or -W for more fine-grained control)
- o Maximum number of packets per output file
- i Maximum number of seconds per output file (Note that the -i and -I flags cause multiple files to be created, each named with an incrementing integer and initial timestamp for each file's content, e.g. output_00000_20170417174516.pcap)

Examples:

```
$ editcap -A '2017-01-16 00:00:00' -B '2017-02-16 00:00:00' infile.pcap
$ editcap -d infile.pcap dedupe.pcap
$ editcap -i 3600 infile.pcap hourly.pcap
```

tshark: Command-line access to nearly all Wireshark features
For all of Wireshark's features, the ability to access them from the command line provides valuable power to the analyst. Whether building repeatable commands into a script, looping over dozens of input files, or performing analysis directly within the shell, **tshark** packs nearly all of Wireshark's features in a command-line interface.

Usage:
\$ tshark -n -r <input file> <options> <filter>

Common command-line parameters:

- n Prevent DNS lookups on IP addresses
- r Read from specified pcap file
- w Write packet data to a file
- Y Specify Wireshark-compatible display filter
- T Specify output mode (Fields, text (default), pdml, etc.)
- e When used with -T fields, specifies a field to include in output tab-separated values (can be used multiple times)
- G Specify glossary to display (protocols, fields, etc.) – shows available capabilities via command line, suitable for grep'ing, etc.

Display filter resources:
See the [Wireshark-filter](#) man page for more command-line details on how to construct display filters.

Examples:

```
$ tshark -n -r infile.pcap -Y "http.host contains \"google\""
$ tshark -n -r infile.pcap -T fields -e ip.src -e http.host
$ tshark -n -r infile.pcap -Y "ssl handshake.certificates"
$ tshark -n -r infile.pcap -w just_certificates.pcap
```

tcpextract: Carve reassembled TCP streams for known header and footer bytes to attempt file reassembly
This is the TCP equivalent to the venerable **foremost** and **scaple**, diskimager, and **carve**. **tcpextract** will reassemble each TCP stream, then search for known header bytes in the stream, writing out matching sub-streams to disk. It is not protocol-aware, so it cannot determine metadata such as filenames and cannot handle protocol content consisting of non-contiguous byte sequences. Notably, **tcpextract** cannot parse SMB traffic, encrypted payload content, or chunked-encoded HTTP traffic. Parsing compressed data requires signatures for the compressed bytes rather than the corresponding plaintext.

Usage:
\$ tcpextract -r <input file> <options>

Common command-line parameters:

- r Read from specified pcap file
- c Configuration (signature) file to use
- o Place output files into specified directory

Signature format:
• file ext(max_size, start_bytes, end_bytes);
Signature examples:

- +if \$file(3000000, x47:x49:x46\x38\x37)x61, \x00\x30\x3b;
- +rp(40000000, xed\xab\xee\xdb);

Example:
\$ tcpextract -f infile.pcap &
\$ rpm-tcpextract.conf &
\$./

grep: Display lines from input text that match a specified regular expression pattern
Searches input text from a file or via STDIN pipes using extremely flexible and age-old regular expressions. Matching lines are displayed, but output can be fine-grained to address specific analytic requirements.

Usage:
\$ grep <options> <pattern> <input file>

Common command-line parameters:

- i Case-insensitive search
- r Recursively process all files within a directory tree
- a Fully search all files as ASCII, even if they appear to contain binary data
- l Only display line names that contain matches instead of the lines on which the match is found
- F Perform the regular expression engine, providing a significant speed benefit
- c Display count of matching lines
- B Display a number of lines before each line that matches the search pattern
- D Display a number of lines after each line that matches the search pattern
- C Display a number of context lines before and after each line that matches the search pattern
- H Display filenames in addition to matching line contents – this is the default with -c
- h Omit filenames from output as displayed with -r
- v Invert match – only shows results that do not match the search pattern – with -l, show files' names in which there is at least one match not matching the search pattern – with -c, show count of non-matching lines

Regular expressions are a dark art of shell commands.

Examples:
\$ grep pastebin.access.log
\$ grep -r google /var/spool/squid/
\$ grep -Fv 192.168.75. syslog-messages
\$ grep -C 5 utmerror.log

NetworkMiner: Protocol-aware object extraction tool that writes files to disk
Object extraction is often a tedious task, but NetworkMiner reliably performs this function for a number of common protocols. File objects are written to disk as they are encountered, while fields (credentials, hosts, etc.) can be exported to CSV format.

Writing files to disk often triggers host-based defenses, so running this utility in an isolated and controlled environment is the most common use model. NetworkMiner is a commercial utility that also provides a free version. The free version is licensed for operational use, not just testing.

zeek-cut: Extract specific fields from Zeek logs
The Zeek NSM creates log files as needed to document observed network traffic. If the tab-separated value (TSV) format is used, the "zeek-cut" utility can extract just the fields of interest.

Usage:
\$ cat <log file> | zeek-cut <options> <fields>

Common command-line parameters:

- u Convert timestamp to human-readable, UTC format
- o Display header blocks at start of output

Identifying fields of interest:
Each different log file type contains various fields, detailed in the header of the file. Inspect the first few lines and identify the one that begins with the string #Fields. The remainder of this line contains the Zeek-specific names for each column of data, which can be extracted with the "zeek-cut" utility. Consult the Zeek NSM documentation for details on each column's meaning.

Examples:
\$ cat files.log | zeek-cut -u ts
\$ cat files.log | zeek-cut -u ts.id.orig_h
\$ cat files.log | zeek-cut -u ts.id.orig_h &
\$ host uri user_agent.info_code

capinfos: Calculate and display high-level summary statistics for an input pcap file
This utility displays summary metadata from one or more source pcap files. Reported metadata includes but is not limited to start/end times, hash values, packet count, and byte count.

Usage:
\$ capinfos <options> <input file 1> <input file 2> <...>

Common command-line parameters:

- A Generate all available statistics
- T Use "table" output format instead of list format

Examples:

```
$ capinfos -A infile.pcap
$ capinfos -A -T infile2.pcap
$ capinfos -A *.pcap
```

nfdump: Process NetFlow data from nfcapd-compatible files on disk
Files created by **nfcapd** (live collector) or **rfpcap** (pcap-to-NetFlow distillation) are read, parsed, and displayed by **nfdump**. Filters include numerous options and calculated fields, and outputs can be customized to unique analysis requirements.

Usage:
\$ nfdump (-R <input directory path> | -r <pcap file>) <options>

Common command-line parameters:

- R Read from the specified single file
- r Recursively read from the specified directory tree
- t Specify time window in which to search (use format: YYYY/MM/DD, hh:mm:ss-MM/DD:hh:mm:ss)
- o Output format to use (line, long, extended, or custom with fmt=<format string>)
- O Output sort ordering (tstart,bytes,packets,more)
- a Aggregate output on source IP-port, destination IP-port, layer 4 protocol
- A Comma-separated custom aggregation fields

Filter syntax:

```
host IP address or FQDN
net Netblock in CIDR notation
proto Layer 4 protocol (tcp, udp, icmp, etc)
as Autonomous System number
Parameters such as host, net, and port can be applied in just one direction with the src or dst modifiers. Primitives can be combined with and, or, or not, and/or can be enforced with parentheses.
```

Filter examples:

- * proto tcp and port 80
- * proto udp and dst host 8.8.8.8
- * src host 1.2.3.4 and (dst net 10.0.0.0/8 or dst net 172.16.0.0/12)
- * src as 32625 (Note: not all collections include ASNs)

Custom output formatting:

```
Format strings for the custom output format option
{ts} !fmt=<format_string>|> custom of format tags, including but not limited to those below:
{ts} Start time
{te} End time
{td} Duration (in seconds)
{sp} Source port (TCP or UDP)
{dp} Destination port (TCP or UDP); formatted as {type}_code for ICMP
{sa} Source IP address and port
{da} Destination IP address and port
{pk} Packet count
{byt} Byte count
{flg} TCP flags (sum total for flow)
{bps} Bits per second (average)
{pps} Packets per second (average)
{app} Bytes per packet (average)

```

Custom aggregation:

```
Records displayed can be aggregated (tallied) on user-specified fields including but not limited to those below:
proto Layer 4 protocol
srcip Source IP address
dstip Destination IP address
srcport TCP or UDP source port
dstport TCP or UDP destination port
srcnet Source netblock in CIDR notation
dstnet Destination netblock in CIDR notation
```

Examples:

```
$ nfdump -r nfpcap.201703271745 &
-o long | proto tcp and port 53'
$ nfdump -R /var/log/netflow/2017/03/ &
-o fnt!${sa} ${da} ${pr} &
-A srcip,dstip,prot
$ nfdump -R /var/log/netflow/2016/ &
-O ststart 'proto tcp and port 4444'
```

tcpflow: Reassemble input packet data to TCP data segments
This utility will perform TCP reassembly, then output each side of the TCP data flows to separate files. This is essentially a scalable, command-line equivalent to Wireshark's "Follow TCP Stream" feature. Additionally, **tcpflow** can perform a variety of decoding and post-processing functions on the resulting flows.

Usage:
\$ tcpflow <options> -r <input file> <pattern> <bpf filter>

Common command-line parameters:

- r Read from specified pcap file
- o Write matching packets to specified pcap file
- i Case-insensitive search
- v Invert match – only show packets that do not match the search pattern
- t Show timestamp from each matching packet

Note: The BPF filter is an optional parameter

Examples:

```
$ ngrep -I infile.pcap 'RETR' 'top and port 21'
$ ngrep -I infile.pcap -i 1337AUTH
```

tcpflow: Reassemble input packet data to TCP data segments
This utility will perform TCP reassembly, then output each side of the TCP data flows to separate files. This is essentially a scalable, command-line equivalent to Wireshark's "Follow TCP Stream" feature. Additionally, **tcpflow** can perform a variety of decoding and post-processing functions on the resulting flows.

Usage:
\$ tcpflow <options> -r <input file> <pattern> <bpf filter>

Common command-line parameters:

- r Read from specified pcap file (can be used multiple times for multiple files)
- l Read from multiple pcap files (with wildcards)
- o Place output files into specified directory

Examples:

```
$ tcpflow -r infile.pcap -o /tmp/output/
$ tcpflow -r -l *.pcap -o /tmp/output/
```

jq: Parse and format JSON data
JSON (JavaScript Object Notation) is a standardized format for key-value pairs and related data structures and is used increasingly for log file content. The "jq" utility provides countless ways to parse and format JSON data.

Usage:
\$ cat <input file> | jq <expression>

Common command-line parameters:

- c Display output in compact format
- r Output raw (unquoted) strings

Notes: Using "...," as the expression will pretty-print the entire input set. UNIX epoch times can be converted to ISO8601 format with the "-r" modifier.

The "select" function can be much slower than using "grep" first and then applying "jq" transforms.

Examples:

```
$ cat file.json | jq '.'
$ cat file.json | jq 'ts: .ts | todate, uid: .uid'
$ cat file.json | jq 'select(.host == "52.0.0.1") | .url'
$ grep -Pzf52.0.0.1 file.json | jq '.url'
$ grep -Pzf52.0.0.1 file.json | jq '.url'
```

calamaris: Generate summary reports from web proxy server log files
The **calamaris** utility performs high-level summary analysis of many different formats of web proxy log files. These reports are broken down by HTTP request methods, second-level domains, client IP addresses, HTTP response codes, and more.

Usage:
\$ cat <input file> | calamaris <options>

Common command-line parameter:

- a Generate all available reports

Examples:

```
$ cat access.log | calamaris -a
$ grep 1.2.3.4 access.log.gz | calamaris -a
$ grep badhost.cc.gz | calamaris -a
```

mergecap: Merge two or more pcap files
When faced with a large number of pcap files, it may be advantageous to merge a subset of them to a single file for streamlined processing. This utility will ensure the packets written to the output file are chronological.

Usage:
\$ mergecap <options> -w <output file> <input file 1> <input file 2> <input file n>

Common command-line parameters:

- w New pcap file to create, containing merged data
- s Number of bytes per packet to retain

Example:
\$ mergecap -w new.pcap infile1.pcap infile2.pcap

Network Traffic Anomalies

HTTP GET vs POST Ratio

How: HTTP proxy logs, NSM logs, HTTP server logs
What: The proportion of observed HTTP requests that use the GET, POST, or other methods.

Why: This ratio establishes a typical activity profile for HTTP traffic. When it skews too far from the normal baseline, it may suggest brute force logins, SQL injection attempts, RAT usage, server feature probing, or other suspicious/malicious activity.

Top-Talking IP Addresses

How: NetFlow
What: The list of hosts responsible for the highest volume of network communications in volume and/or connection count. Calculate this on a rolling daily/weekly/monthly/annual basis to account for periodic shifts in traffic patterns.
Why: Unusually large spikes in traffic may suggest exfiltration activity, while spikes in connection attempts may suggest C2 activity.

HTTP User-Agent

How: HTTP proxy logs, NSM logs, HTTP server logs
What: The HTTP User-Agent generally identifies the software responsible for issuing an HTTP request. This can be useful to profile software operating within the environment.
Why: This is an invaluable identifier to profile activity within the environment. It can profile which web browser titles, versions, and extensions are in use. Recently, desktop and mobile applications use unique User-Agent strings as well. Knowing the "normal" string presents causes outliers to stand out, which may highlight suspicious activity. However, this is an arbitrary and optional header, so be skeptical of behavior that suggests forgery – such as rapid change for a given IP address, significant increase in the number of observed User-Agent strings, etc.

Top DNS Domains Queried

How: Passive DNS logs, DNS server-side query logs, NSM logs
What: The most frequently queried second-level domains (e.g. "example.com" or "example.co.uk") based on internal clients' request activity. The top 1000 domains on a rolling daily basis may be a good starting point, but this number should be adjusted to local requirements.
Why: In general, the behaviors of a given environment don't drastically change on a day-to-day basis. Therefore, the top 500-700 domains queried on any given day should not differ too much from the top 1000 from the previous day. The difference in count allows for natural ebb and flow of daily behavior. Any domain that rockets to the top of the list may suggest an event that requires attention, such as a new phishing campaign, C2 domain, or other anomaly.

HTTP Return Code Ratio

How: HTTP proxy logs, NSM logs, HTTP server logs
What: The return code is a three-digit integer that helps to indicate "what happened" on the server answering a request. These are grouped into "families" by hundreds: 100s = informational, 200s = success, 300s = redirection, 400s = client-side error, 500s = server-side error.
Why: Knowing what happened at the server end of the transaction can be extremely useful in characterizing HTTP activity. A spike in 400-series codes could indicate reconnaissance or scanning activity, while an unusually high number of 500-series codes could indicate failed login or SQL injection attempts. As with other observations, knowing the typically-observed ratios of these values can help to identify anomalous trends that require further investigation.

Newly-Observed/Newly-Registered Domains

How: Passive DNS logs, DNS server-side query logs, NSM logs
What: Any domain that has never previously been queried from within the environment, according to its WHOIS "Date Registered."
Why: The first time a domain is queried in a given environment may indicate a new or highly-focused targeting operation. Brand new domains are often associated with malicious activity, given that attackers generally require a dynamic infrastructure for their operations.

External Infrastructure Usage Attempts

How: NetFlow, Firewall logs, NSM logs
What: Although best practice is to restrict outbound communications by default and approve necessary services and connections by exception, this is often not the case – perimeter are still notoriously porous in the outbound direction. Even in a properly-constrained environment, these attacks should create artifacts of the failed connection attempts.
Why: By identifying internal clients that attempt to or succeed in using external services, it is possible to quickly collect a list of endpoints that exhibit anomalous behavior. These may include connections to external DNS servers rather than internal resolvers, HTTP connection attempts that seek to bypass proxy servers, connections to VPN providers, raw socket connections to unusual ports, and more.

Typical Port and Protocol Usage

How: NetFlow
What: The list of ports and corresponding protocols that account for the most communication in terms of volume and/or connection count. Calculate this on a daily/weekly/monthly/annual basis to account for periodic shifts in traffic patterns.
Why: Similar to the purpose for tracking top-talking IP addresses, knowing the typical port and protocol usage enables quick identification of anomalies that should be further explored for potential suspicious activity.

DNS TTL Values and RR Counts

How: Passive DNS logs, NSM logs
What: TTL refers to the number of seconds that a caching DNS server should retain a given record. The number of Resource Records in a given DNS packet is noted in the RR count field.
Why: Very short TTLs may suggest fast-flux DNS or potential tunneling behavior. A high RR count could indicate large-scale load balancing associated with fast-flux or similar elastic architectures. While these behaviors can suggest suspicious behavior, they are also commonly seen with benign network activity such as content delivery networks, round robin DNS-based load balancing, and similar architectures.

Autonomous System Communications

How: NetFlow, NSM logs
What: Autonomous System Numbers (ASNs) are numerical 'handles' assigned to netblocks owners such as ISPs, datacenters, and other service providers. These can suggest Internet "neighborhoods" to characterize network traffic based on more than IP address or CIDR blocks.
Why: Certain ASNs are often more prominently associated with malicious activity than others. Reputation databases can be useful in determining these. Even without an intelligence overlay, identifying the ASNs with which systems in the environment communicate is a useful baseline metric that can easily identify communications with unusual ASNs that require further attention.

Periodic Traffic Volume Metrics

How: NetFlow
What: Maintaining traffic metrics on time-of-day, day-of-week, day-of-month, and similar bases.
Why: These will identify normative traffic patterns, making deviations easier to spot and investigate. A sudden spike of traffic or connections during an overnight or weekend period (when there is typically little or no traffic) would be a clear anomaly of concern.

Arkime

Arkime is a full-packet ingestion and indexing platform. It reads a live network data stream or existing pcap files, then extracts data from known protocol fields to store in an Elasticsearch index, associating the client- and server-sourced directions of a connection for easy analysis. Arkime separates full-packet data and SPI data, allowing different storage allocation and retention policies. The user can export a subset of traffic in pcap format, making it a valuable addition to the workflow, since any pcap-aware tool can be used on the derived data.

Loading Data to Arkime

Arkime can load network traffic from existing pcap files (DFIR Model) or a live network interface (Security Operations Model).

DFIR Model

Load pcap files with the "moloch-capture"